

Solving the N + 1 Query problem

For Microservices Without any framework

Agenda

- What is the N + 1 query problem?
- Why it is harder to solve in a microservice architecture?
- How to build a user-friendly API to solve the problem?

Assumptions

- The talk is really about building an API using Java 8+ functional patterns
- Not about databases, JPA or Hibernate
- Mostly not about high level microservice architectures
- Very code oriented focusing on core java

What is the N + 1 query problem?

Example: *Online Store*

Get a list of all items purchased by some customers

```
Select * from customer where...
```

What is the N + 1 query problem?

Assuming we have 10 customers, for each 10 customers c :

```
Select * from order_item i  
      where i.customer_id = c.customer_id
```

What is the N + 1 query problem?

How many queries issued to the database?

What is the N + 1 query problem?

Answer: 11 (10 + 1)

- 1 query to `customer`
- 10 queries to `order_item`

What is the N + 1 query problem?

This is linear... until... it's not anymore!

- 100 customers... 101 queries
- 500 customers... 501 queries
- 1000 customers... 0 query ;-)

So what should happen instead?

- 10, 100, 500 customers... doesn't matter
- Only 1 query to `order_item`
- Join each java Customer entity with each associated OrderItem

ORMs to the rescue?

- JPA query with Hibernate for example (pseudo-code):

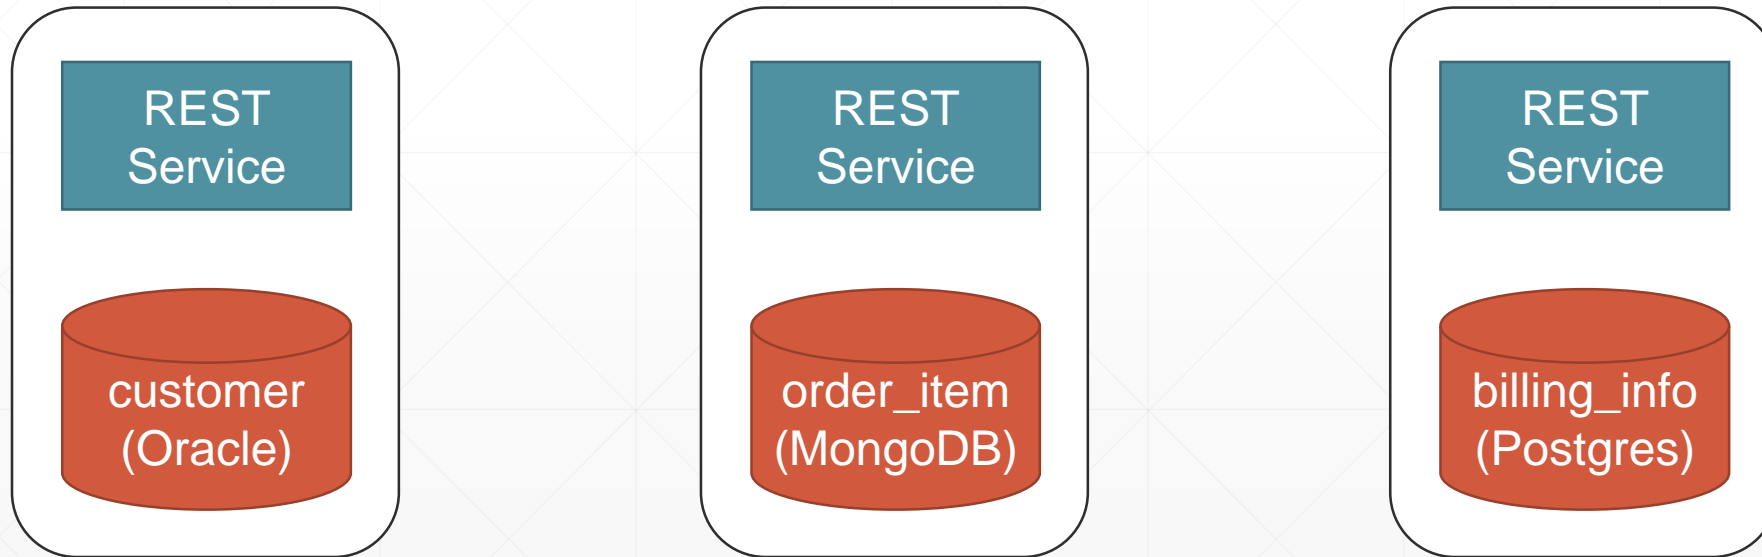
```
Select c from Customer c  
    JOIN FETCH c.orderItems
```

Will execute behind the scene a query looking like:

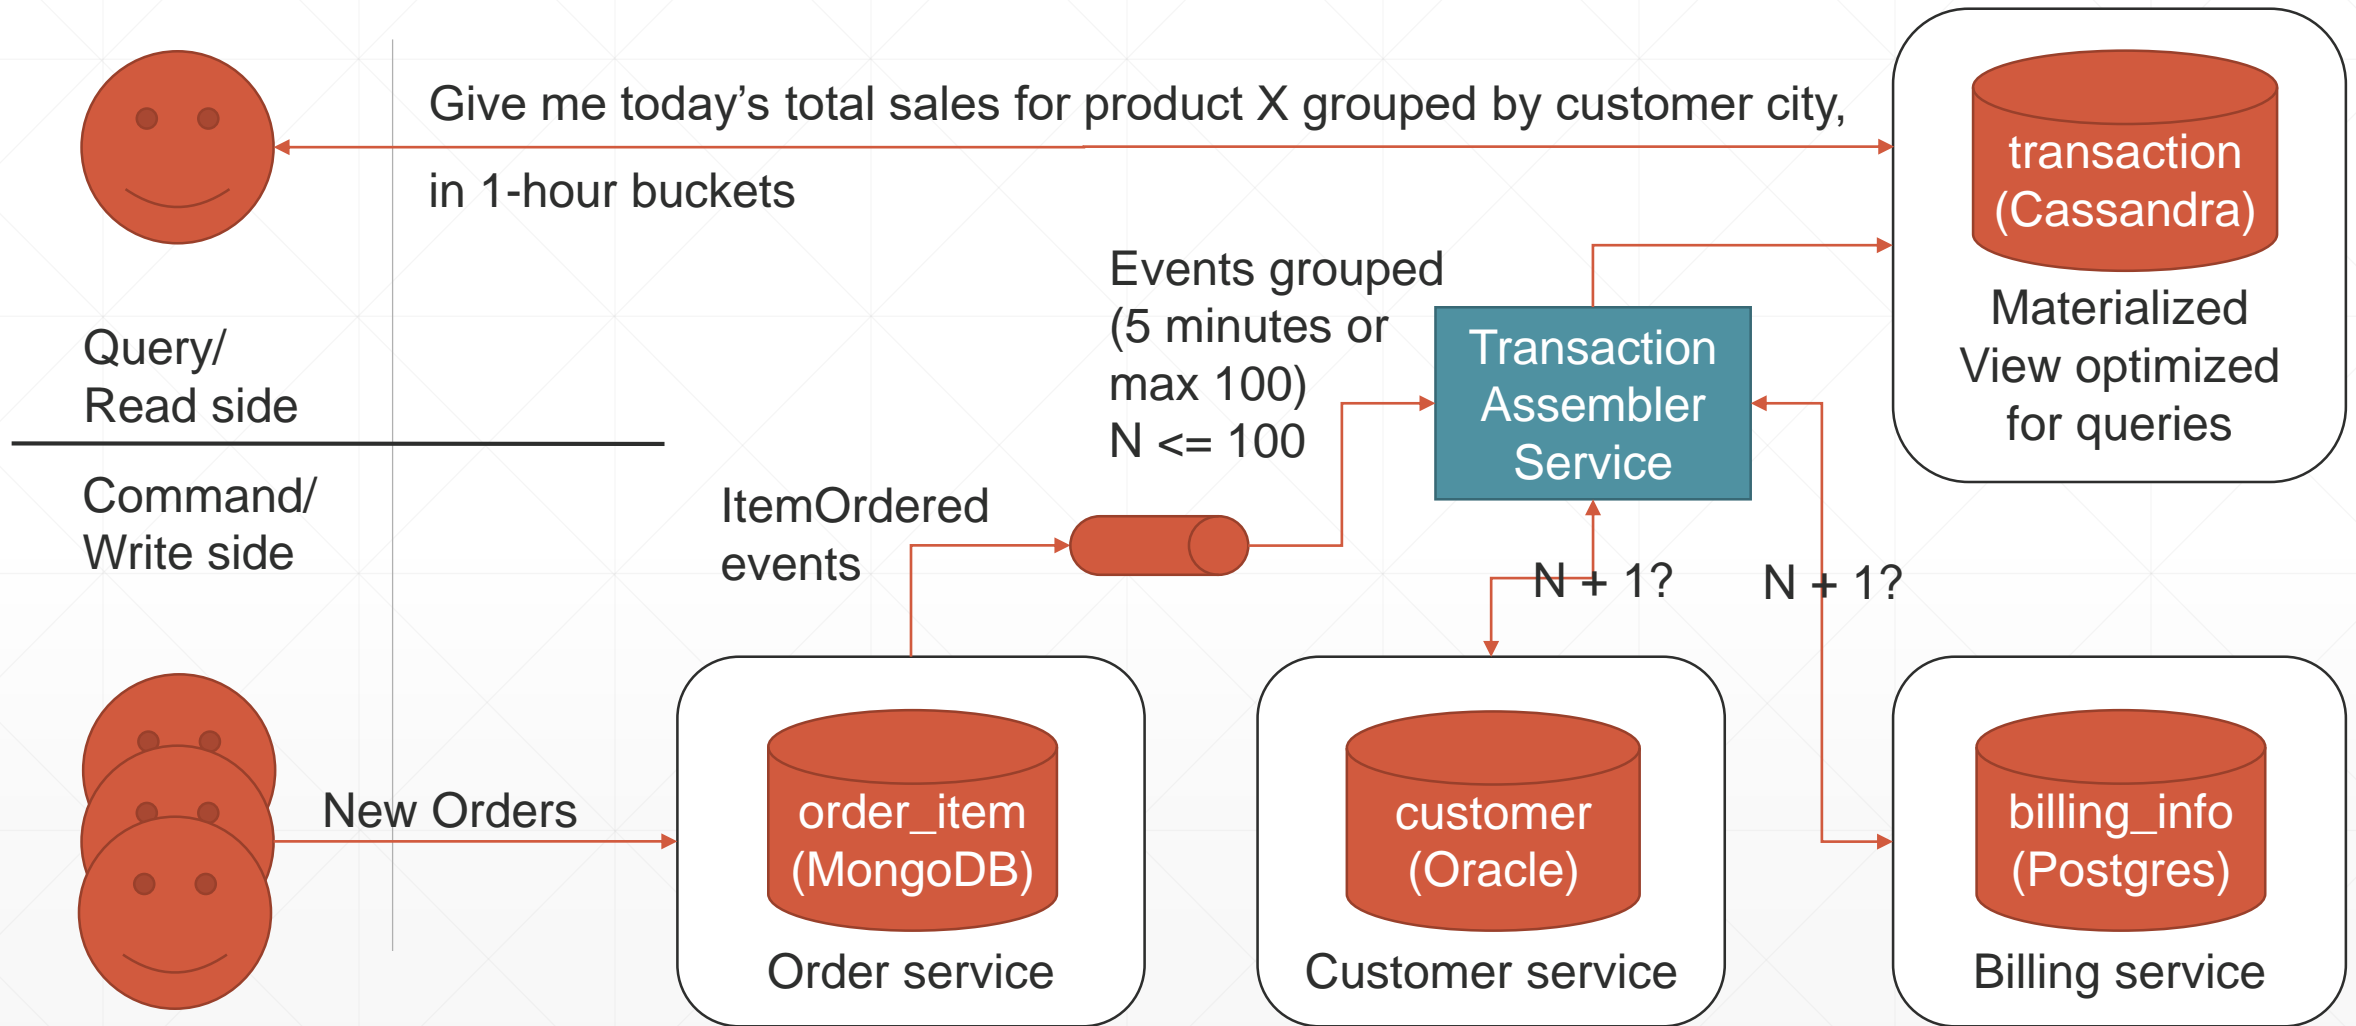
```
select ... from customer c0  
    inner join order_item o1  
    on c0.customer_id = o1.customer_id
```

But in a Microservice Environment...

- ORMs works with a single database
- Now we need to aggregate disconnected sources of data...

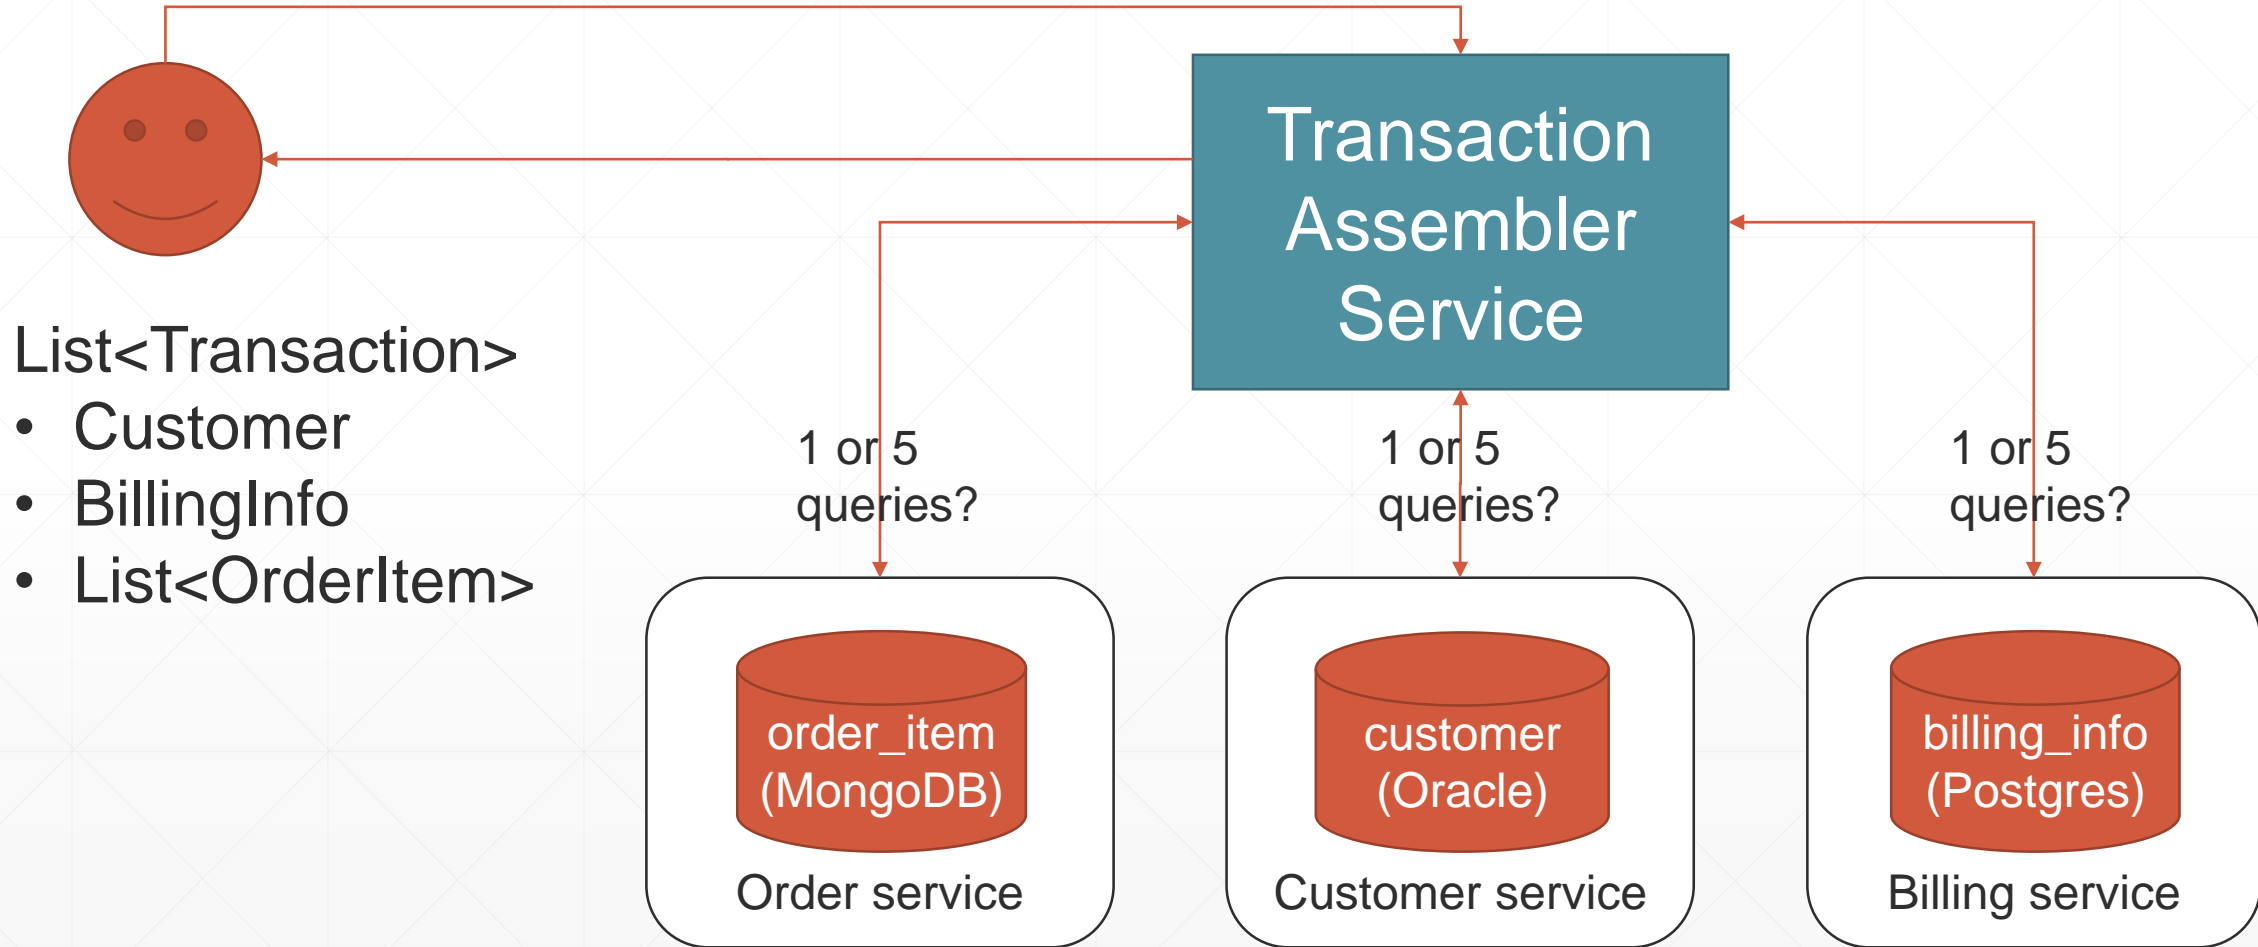


Fictional ES/CQRS Microservice Scenario



A Simpler Example

Give me all transactions for customer id 1, 2, 3, 4, 5



So how do we implement that Box?

- Without triggering the N + 1 Query Problem

Transaction
Assembler
Service

Assuming...

```
@Data
@AllArgsConstructor
public class Customer {
    private final Long customerId;
    private final String name;
}
```

And...

```
@Data
@AllArgsConstructor
public class BillingInfo {
    private final Long billingInfoId
    private final Long customerId;
    private final String creditCardNumber;
}
```


And...

@Data

@AllArgsConstructor

```
public class OrderItem {  
    private final Long orderId  
    private final Long customerId;  
    private final String orderDescription;  
    private final Double price;  
}
```

And Our Aggregate...

```
@Data
@AllArgsConstructor
public class Transaction {
    private final Customer customer;
    private final BillingInfo billingInfo;
    private final List<OrderItem> orderItems;
}
```

Now our disparate data sources...

Spring Data JPA Repository on Oracle Database

```
@Query("select c from Customer c where c.customerId in :ids")  
List<Customer> getCustomers(List<Long> customerIds);
```

And...

REST HTTP or Binary Protocol Client to BillingInfo Service

```
List<BillingInfo> getBillingInfo(List<Long> customerIds);
```

Notice the method signature?

And...

Spring Data MongoDB Repository

```
@Query("{ 'customerId' : { $in: ?0 } }")  
List<OrderItem> getAllOrders(List<Long> customerIds);
```

Notice again the method signature?

This is the API we want to build

```
Assembler<Customer, Flux<Transaction>> assembler =  
  
  assemblerOf(Transaction.class)  
    .withIdExtractor(Customer::getCustomerId)  
    .withAssemblerRules(  
      oneToOne(this::getBillingInfo, BillingInfo::getCustomerId),  
      oneToMany(this::getAllOrders, OrderItem::getCustomerId),  
      Transaction::new)  
    .using(fluxAdapter()); // or Stream adapter, RxJava adapter, etc.
```

Usage example with Spring Cloud Stream (Kafka) + Project Reactor

```
private Assembler<Customer, Flux<Transaction>> assembler = ...

@StreamListener
@Output(Processor.OUTPUT)
public Flux<Transaction> receive(
    @Input(Processor.INPUT) Flux<Customer> customerFlux) {

    return customerFlux.bufferTimeout(100, ofMinutes(5))
        .flatMapSequential(assembler::assemble);
}
```

How do we build such an API?

1. Implement oneToOne and oneToMany semantics for sub-queries
2. Join/Aggregation logic
3. Decouple Join logic from execution engines
 - Java 8 Stream
 - Project Reactor
 - RxJava
 - Akka Stream


1. Implement oneToOne and oneToMany semantics

```
assemblerOf(Transaction.class)  
  .withIdExtractor(Customer::getCustomerId)  
  .withAssemblerRules(  
    oneToOne(this::getBillingInfo, BillingInfo::getCustomerId),  
    oneToMany(this::getAllOrders, OrderItem::getCustomerId),  
    Transaction::new)  
  .using(fluxAdapter());
```

Let's start with a simple utility method...

```
Map<Long, BillingInfo> billingInfoMap =  
    queryOneToOne(of(1L, 2L, 3L),  
        this::getBillingInfo, // REST HTTP call  
        BillingInfo::getCustomerId);
```

List<BillingInfo>		
B1(1, ...)	B2(2, ...)	B3(3, ...)




Map<Long, BillingInfo>	
CustomerID	BillingInfo
1	B1
2	B2
3	B3

Let's start with a simple utility method...

```
Map<Long, List<OrderItem>> orderItemMap =  
    queryOneToMany(of(1L, 2L, 3L),  
        this::getAllOrders, // Spring Data MongoDB  
        OrderItem::getCustomerId,  
        ArrayList::new);
```

List<OrderItem>				
O1(1, ...)	O2(1, ...)	O3(1, ...)	O4(2, ...)	O5(2, ...)



Map<Long, List<OrderItem>>	
CustomerId	OrderItems
1	{ O1, O2, O3 }
2	{ O4, O5 }
3	{ }

queryOneToOne vs. queryOneToMany

- Both queryOneToOne and queryOneToMany are very similar
- They both retrieve list of data for a data source
- They both group the data received (by ID)
- They only differ in how they group the data

Strategy pattern with Java 8 Stream collectors

```
static <ID, R, IDC extends Collection<ID>, RC extends  
Collection<R>>
```

```
Map<ID, R> queryOneToOne(IDC ids,  
    Function<IDC, RC> queryFunction,  
    Function<R, ID> idExtractor) {  
  
    return query(ids, queryFunction,  
        toMap(idExtractor, identity()));  
}
```

Strategy pattern with Java 8 Stream collectors

```
Map<Long, BillingInfo> queryOneToOne(List<Long> customerIds,  
    Function<List<Long>, List<BillingInfo>> queryFunction,  
    Function<BillingInfo, Long> idExtractor) {  
  
    return query(customerIds, queryFunction,  
        toMap(idExtractor, identity()));  
}
```

Strategy pattern with Java 8 Stream collectors

```
Map<ID, RC> queryOneToMany(IDC ids,  
    Function<IDC, RC> queryFunction,  
    Function<R, ID> idExtractor,  
    Supplier<RC> collectionFactory) {  
  
    return query(ids, queryFunction,  
        groupingBy(idExtractor,  
            toCollection(collectionFactory))) ;  
}
```

Strategy pattern with Java 8 Stream collectors

```
Map<Long, List<OrderItem>> queryOneToMany(List<Long> customerIds,  
    Function<List<Long>, List<OrderItem>> queryFunction,  
    Function<OrderItem, Long> idExtractor,  
    Supplier<List<OrderItem>> collectionFactory) {  
  
    return query(customerIds, queryFunction,  
        groupingBy(idExtractor,  
            toCollection(collectionFactory))) ;  
}
```


Strategy pattern with Java 8 Stream collectors

Our generic query() method:

```
Map<ID, V> query(IDC ids,  
                 Function<IDC, RC> queryFunction,  
                 Collector<R, ?, Map<ID, V>> mapCollector) {  
  
    return queryFunction.apply(ids).stream()  
        .collect(mapCollector);  
}
```

$V = R \parallel RC$ (e.g. `List<R>`)

From queryOneToXXX to oneToXXX

Function returning function: From eager to lazy execution

```
Mapper<ID, R> oneToOne (
    Function<IDC, RC> queryFunction,
    Function<R, ID> idExtractor) {

    return entityIds -> queryOneToOne(
        (IDC) entityIds, queryFunction, idExtractor);
}
```

Function returning function: From eager to lazy execution

```
@FunctionalInterface
public interface Mapper<ID, R> {

    Map<ID, R> apply(Iterable<ID> entityIds);

}
```

- Like `java.util.Function` with predefined types
- Equivalent to `Function<Iterable<ID>, Map<ID, R>>`

Function returning function: From eager to lazy execution

```
Mapper<Long, BillingInfo> billingInfoMapper =  
    oneToOne(this::getBillingInfo, BillingInfo::getCustomerId);
```

- Move `billingInfoMapper` around, pass it to other methods...
- Later invoke it...

```
Map<Long, BillingInfo> map =  
    billingInfoMapper.apply(of(1L, 2L, 3L));
```

Function returning function: From eager to lazy execution

You can guess the implementation of `oneToMany`...

Function returning function: From eager to lazy execution

```
Mapper<ID, RC> oneToMany(  
    Function<IDC, RC> queryFunction,  
    Function<R, ID> idExtractor,  
    Supplier<RC> collectionFactory) {  
  
    return entityIds -> queryOneToMany(  
        (IDC) entityIds, queryFunction, idExtractor,  
        collectionFactory);  
}
```

2. Implement Join/Aggregate Logic (part 1)

```
assemblerOf(Transaction.class)  
  .withIdExtractor(Customer::getCustomerId)  
  .withAssemblerRules(  
    oneToOne(this::getBillingInfo, BillingInfo::getCustomerId),  
    oneToMany(this::getAllOrders, OrderItem::getCustomerId),  
    Transaction::new)  
  .using(fluxAdapter());
```


We could have 3, 5, 10 overloaded methods...

```
withAssemblerRules(  
    Mapper<ID, E1> mapper,  
    BiFunction<T, E1, R> assemblerFunction)
```

```
withAssemblerRules(  
    Mapper<ID, E1> mapper1,  
    Mapper<ID, E2> mapper2,  
    Function3<T, E1, E2, R> assemblerFunction)
```

But wouldn't it be so easier to just use...

```
withAssemblerRules(  
    List<Mapper<ID, ?>> mappers,  
    BiFunction<T, Object[], R> aggregationFunction);
```




... and reuse the same code for an arbitrary number of parameters instead?

**Sure, except for the fact that we are losing
type information...**

We want the compiler to catch this...

```
withAssemblerRules(  
  List.of(  
    oneToMany(this::getAllOrders, OrderItem::getCustomerId),  
    oneToOne(this::getBillingInfo, BillingInfo::getCustomerId)  
  ),  
  (Customer c, Object[] subQueryResults) ->  
    new Transaction(c, (BillingInfo) subQueryResults[0],  
                    (List<OrderItem>) subQueryResults[1])  
)
```



This smells ClassCastException using a generic solution...

So how to get from...


```
BiFunction<T, Object[], R> aggregationFunction
```

Back to...

```
BiFunction<T, E1, R> assemblerFunction)
```

```
Function3<T, E1, E2, R> assemblerFunction)
```

Double Sided Funnel Pattern (You won't find it on the Internet 😊)



- Type information available

- Type information loss

- Type information restored

Mapper<ID, E1> mapper1

Mapper<ID, E1> mapper1
Mapper<ID, E2> mapper2

Mapper<ID, E1> mapper1
Mapper<ID, E2> mapper2
Mapper<ID, E3> mapper3

List<Mapper<ID, ?>> mappers

BiFunction<T, Object[], R>
aggregationFunction

BiFunction<T, E1, R> assemblerFunction

Function3<T, E1, E2, R> assemblerFunction

Function4<T, E1, E2, E3, R> assemblerFunction

- Input parameters expanded

- Input parameters collapsed in List
- Output values collapsed in Object[]

- Output values expanded

Assuming this...

```
Function3<T, E1, E2, R> assemblerFunction =
```

```
(Customer c, BillingInfo b, List<OrderItem> o) -> new Transaction(c, b, o))
```

... with our generic method called by our 2, 5, 10 overloaded methods
of withAssemblerRules

```
withAssemblerRules(  
    List<Mapper<ID, ?>> mappers,  
    BiFunction<T, Object[], R> aggregationFunction);
```


Double Sided Funnel Pattern with 2 Parameters

```
<E1, E2> AssembleUsingBuilder<T, ID, R> withAssemblerRules(  
    Mapper<ID, E1> mapper1,  
    Mapper<ID, E2> mapper2,  
    Function3<T, E1, E2, R> assemblerFunction) {  
  
    return withAssemblerRules(List.of(mapper1, mapper2),  
        (topLevelEntity, mapperResults) ->  
            assemblerFunction.apply(topLevelEntity,  
                                    (E1) mapperResults[0],  
                                    (E2) mapperResults[1])) ;  
}
```

Double Sided Funnel Pattern with 3 Parameters

```
<E1, E2, E3> AssembleUsingBuilder<T, ID, R> withAssemblerRules(  
    Mapper<ID, E1> mapper1,  
    Mapper<ID, E2> mapper2,  
    Mapper<ID, E3> mapper3,  
    Function4<T, E1, E2, E3, R> assemblerFunction) {  
  
    return withAssemblerRules(List.of(mapper1, mapper2, mapper3),  
        (topLevelEntity, mapperResults) ->  
            assemblerFunction.apply(topLevelEntity,  
                                    (E1) mapperResults[0],  
                                    (E2) mapperResults[1],  
                                    (E3) mapperResults[2]));  
}
```

2. Implement Join/Aggregate Logic (part 2)

```
assemblerOf(Transaction.class)  
  .withIdExtractor(Customer::getCustomerId)  
  .withAssemblerRules(  
    oneToOne(this::getBillingInfo, BillingInfo::getCustomerId),  
    oneToMany(this::getAllOrders, OrderItem::getCustomerId),  
    Transaction::new)  
  .using(fluxAdapter());
```

Join/Aggregate function

For 3 customers (or 20, 40, etc.):

- ALWAYS only 2 network calls
 - getBillingInfo
 - getAllOrders
- For each **customer**
 - For each **Map<Long, ?>**
 - **Map.get(customerId)**
- **3** x **2** = 6 in-memory iterations

List<Customer>
C1
C2
C3

From oneToOne(...)	
Map<Long, BillingInfo>	
CustomerID	BillingInfo
1	B1
2	B2
3	B3
From oneToMany(...)	
Map<Long, List<OrderItem>>	
CustomerID	OrderItems
1	{ O1, O2, O3 }
2	{ O4, O5 }
3	{ }

Join/Aggregate function

- For each **customer**
 - For each `Map<Long, ?>`
 - `Map.get(customerId)`

This is the output when called for each 3 customers:

```
aggregationFunction.apply( C1, [ B1, { O1, O2, O3 } ] )  
aggregationFunction.apply( C2, [ B2, { O4, O5 } ] );  
aggregationFunction.apply( C3, [ B3, { } ] );
```

```
BiFunction<T, List<Map<ID, ?>>, R> joinMapperResultsFunction =  
    (topLevelEntity, listOfMapperResults) ->  
        aggregationFunction.apply(topLevelEntity,  
            listOfMapperResults.stream()  
                .map( mapperResult ->  
                    mapperResult.get( idExtractor.apply(topLevelEntity) ) )  
                .toArray() );
```

Join/Aggregate function

So we have our Join/Aggregate logic wrapped in a reusable function...

But nothing is executed yet, so far we are just writing the recipe...

It's Execution Time!!!

```
RC assemble(C topLevelEntities, // List<Customer>  
    Function<T, ID> idExtractor, // Customer::getCustomerId  
    List<Mapper<ID, ?>> subQueryMappers,  
    BiFunction<T, Object[], R> aggregationFunction,  
    AssemblerAdapter<ID, R, RC> assemblerAdapter)
```

Step 1

- Extract customer ids from our list of customers:

```
List<ID> ids = topLevelEntities.stream()  
    .map(idExtractor) // Customer::getCustomerId  
    .collect(toList());
```


Step 2

- Convert our mappers into `Supplier<Map<ID, ?>>`
- i.e. we transform 1 parameter function into 0 parameter function by capturing **entityIDs** in closure
 - Like we did with `queryOneToXXX` earlier

```
Stream<Supplier<Map<ID, ?>>> mapperSourceSuppliers =  
    subQueryMappers.stream()  
        .map (mapper -> () -> mapper.apply(ids));
```

Step 3

- For each **customer**:
 - Reuse our `joinMapperResultFunction`

```
Function<List<Map<ID, ?>>, Stream<R>> aggregateStreamBuilder =  
    mapperResults -> topLevelEntities.stream()  
        .map(topLevelEntity ->  
            joinMapperResultsFunction.apply(topLevelEntity,  
                                             mapperResults));
```

`<R> = <Transaction>`

Step 4

- Let pluggable execution engine do the work
 - Java 8 Steams
 - CompletableFuture
 - Project Reactor
 - RxJava
 - Akka Stream

```
return assemblerAdapter.convertMapperSources (  
    mapperSourceSuppliers, aggregateStreamBuilder) ;
```

3. Decoupling Join/Aggregation Logic from Execution Model

```
assemblerOf(Transaction.class)  
  .withIdExtractor(Customer::getCustomerId)  
  .withAssemblerRules(  
    oneToOne(this::getBillingInfo, BillingInfo::getCustomerId),  
    oneToMany(this::getAllOrders, OrderItem::getCustomerId),  
    Transaction::new)  
  .using(fluxAdapter());
```

Decoupling Join/Aggregate Logic from Execution Model

```
/**
 * @param <ID> e.g. {@code <Long>}
 * @param <R> e.g. {@code <Transaction>}
 * @param <RC> e.g. {@code Stream<Transaction>} or {@code Flux<Transaction>}
 */
@FunctionalInterface
public interface AssemblerAdapter<ID, R, RC> {

    RC convertMapperSources (
        Stream<Supplier<Map<ID, ?>>> mapperSourceSuppliers,
        Function<List<Map<ID, ?>>, Stream<R>> aggregateStreamBuilder);
}
```

Adapter Responsibilities...

- Convert `Supplier<Map<ID, ?>>` to adapter specific type
- Trigger the execution of the suppliers
 - Will perform network calls (e.g. `getBillingInfo()`, `getAllOrders()`)
- Pass the returned `List<Map<ID, ?>>` to `aggregateStreamBuilder`
- Done!

Java 8 Stream Adapter

```
@Override
public Stream<R> convertMapperSources (
    Stream<Supplier<Map<ID, ?>>> mapperSourceSuppliers,
    Function<List<Map<ID, ?>>, Stream<R>> aggregateStreamBuilder) {

    List<Map<ID, ?>> mappers = mapperSourceSuppliers
        .map(Supplier::get)
        .collect(toList());

    return aggregateStreamBuilder.apply(mappers);
}
```

Flux Adapter (Project Reactor)

```
@Override
public Flux<R> convertMapperSources(
    Stream<Supplier<Map<ID, ?>>> mapperSourceSuppliers,
    Function<List<Map<ID, ?>>, Stream<R>> aggregateStreamBuilder) {

    List<Publisher<Map<ID, ?>>> publishers = mapperSourceSuppliers
        .map(this::toPublisher)
        .collect(toList());

    return Flux.zip(publishers,
        mapperResults -> aggregateStreamBuilder.apply(Stream.of(mapperResults)
            .map(mapResult -> (Map<ID, ?>) mapResult)
            .collect(toList()))
        .flatMap(Flux::fromStream);
}

private Publisher<Map<ID, ?>> toPublisher(Supplier<Map<ID, ?>> mapperSource) {
    return fromSupplier(mapperSource).subscribeOn(parallel());
}
```


Remember This?

```
private Assembler<Customer, Flux<Transaction>> assembler = ...

@StreamListener
@Output(Processor.OUTPUT)
public Flux<Transaction> receive(
    @Input(Processor.INPUT) Flux<Customer> customerFlux) {

    return customerFlux.bufferTimeout(100, ofMinutes(5))
        .flatMapSequential(assembler::assemble);
```

Thank You!

<http://bit.ly/javatechw>

